



Anamika Bhardwaj

PHONE: +91-8360734072

EMAIL: anamikabhardwaj06@gmail.com

LINKEDIN: [\[Anamika Bhardwaj | LinkedIn\]](#)

Diligent, self-motivated, and results-driven AI Engineer with strong experience in building scalable Generative AI and agent-based systems. Seeking to contribute in a pragmatic and impactful manner within an organization where I can leverage my technical expertise, continuously enhance my skills, and drive business-aligned AI solutions with integrity and enthusiasm.

TECHNICAL SKILLS: -

AI & Machine Learning

- Large Language Models (LLMs): Transformers-based LLMs, Model-agnostic LLM integration using LiteLLM, Multi-turn Conversational AI, Memory & Context Management, Text Classification, Intent Detection, Semantic Search
- Prompt Engineering: ReAct, Tool Calling, System Prompt Design, Multi-step Reasoning.
- Agentic Frameworks and Architectures: LangChain, LangGraph, AutoGen, CrewAI, Strands, Tools-driven agents, State-driven & Graph-based execution, Agent-as-a-Service architectures
- Retrieval-Augmented Generation (RAG): Vector Embeddings, Document Processing, Chunking Techniques, Similarity Search, VectorDB(Chroma, FAISS, Pinecone), Hybrid Retrieval (Structured + Unstructured Data)
- LLM Evaluation, Safety & Observability: Agent & prompt evaluation strategies, Adversarial & Edge case Testing, Logging, Timing, Tracing & Debugging for LLM pipelines
- APIs & AI Services: FastAPI for building agent services, MCP Servers, API testing & validation using Postman.
- Machine Learning: TensorFlow, Exploratory Data Analysis, Clustering, Similarity Search, Customer Analytics, Model Training

Data Engineering & Cloud

- Programming: Python, SQL (including Spark SQL, PySpark), Shell Scripting
- Data Engineering: ETL Pipeline Design, Data Processing, Azure Synapse, Data Quality Systems
- Databases: Vector Databases, Teradata, PostgreSQL, Data Warehousing, Dimensional Modeling
- Cloud Platforms: AWS, Azure Synapse, Databricks, Cloud Storage Solutions
- CI/CD Pipelines: Git, GitHub Actions, Jenkins

Analytics, Architecture & Leadership

- AI System and Agent Architecture Design, KPI Definition for AI Products, Business-Focused GenAI Solutions
- Technical Documentation (Architecture, APIs, Design Docs, Executive Reporting)
- Cross-team Collaboration, Training & Mentorship
- AI Ethics, Trust & Safety Considerations

WORK EXPERIENCE & RESPONSIBILITY UNDERTAKEN: -

Senior AI Engineer | Omnicom Media Group

(August 2025 – Present)

- Developed a Python Code Review Multi-Agent System that performs syntax validation, static analysis, best-practice checks, and context-aware code feedback using ReAct agents and stateful execution nodes.
- Implemented model-agnostic LLM integration using LiteLLM, allowing seamless switching and routing across OpenAI, LLaMA (Ollama), and other providers without application-level changes.
- Architected and delivered production-grade Agentic AI systems using LangChain and LangGraph, leveraging state-driven graphs, tool-based reasoning (ReAct), and multi-agent orchestration.
- Built Agent-as-a-Service platforms, exposing LLM agents as FastAPI endpoints, enabling modular reuse across internal products and workflows.
- Developed RAG based system that searches a FAISS vector Store for relevant PDF content and generates context-aware answers using Ollama LLaMa 3.1, with source citation for transparency.
- Implemented a Product Q&A Agent using LangGraph + LangChain, featuring: Tool-driven reasoning for pricing and feature retrieval, Hybrid document + structured data search, Multi-turn conversational memory with checkpointing, Robust fallback and error-handling mechanisms.
- Led the Core Sub Agent architecture redesign, transitioning from a hardcoded, US-only system to a region-aware, time-aware, LLM-driven extraction and semantic search pipeline.
- Created LLM-powered extraction tools for keyword detection, region identification, and time-period normalization, enabling reusable, composable agent workflows

- Implemented observability and evaluation of pipelines for agents, including performance logging, deterministic execution checks, and edge-case handling.
- Used Postman for validating agent endpoints, testing tool APIs, and ensuring contract consistency across services.
- Partnered with product, data science, and QA teams to align GenAI systems with business requirements, trust & safety standards, and platform scalability goals.

Senior Data Engineering & AI/ML Analyst | Accenture

(May 2021 – July 2025)

TELSTRA

- Spearheaded the design and implementation of an agentic AI framework for Business Intelligence (AI for BI) using Azure OpenAI services, enabling dynamic NL-to-SQL query generation and automated response evaluation workflows.
- Developed and integrated multi-agent systems, including a natural language query-to-SQL generation agent and a response evaluation agent, to streamline analytics interaction within Power BI Copilot and PostgreSQL backend infrastructure.
- Collaborated cross-functionally to enhance system reliability and response quality, establishing custom LLM evaluation metrics for relevance, accuracy, and completeness to ensure actionable BI insights.

NAB (National Australian Bank)

- Led end-to-end solution design and implementation of multiple GenAI initiatives, developing architecture and technical approaches for each use case
- Architected and implemented synthetic customer data generation system using GPT-4, creating detailed personas for Gen-Z and Mass Affluent segments to enable personalized marketing strategies
- Designed and developed an LLM-powered email content restructuring system, including solution architecture, prompt engineering, and HTML parsing components
- Created comprehensive brand feedback system architecture, implementing review mechanisms and analysis pipelines for maintaining brand consistency
- Engineered automated SEM headline generation pipeline, incorporating historical performance data and brand guidelines
- Designed notification systems for both customer-facing and internal communications, ensuring scalability and brand alignment

TERADATA

- Architected and implemented banking chatbot using RAG framework, designing complete data flow and retrieval systems
- Developed scalable PDF embedding system with VectorDB integration for efficient document retrieval
- Led telecom customer churn analysis project, designing data frameworks for advanced analytics
- Innovated solutions for large document processing, optimizing embedding generation for 100+ page documents

AMAZON NEO & TDLC

- Developed LLM-based system for automated Data Quality rule generation
- Architected ETL code generation system using LLM and prompt engineering
- Led 20-member team in developing and curating high-quality training datasets for language models
- Conducted comprehensive chatbot interaction analysis for performance optimization

DIAGEO

- Designed and implemented data models (Logical, Physical, Semantic, Conceptual) for KPI tracking
- Developed and tested SQL views for business intelligence requirements

ADM (ARCHER DANIELS MIDLAND)

- Led source system integration from multiple platforms (SAP, JDE, SAGE, SPECNET, Oil-Soft)
- Designed and implemented E2E data pipelines in Azure Synapse
- Performed comprehensive data validation and analysis

Data Management Associate | Uplers Pvt Ltd

(Aug 2020 - May 2021)

- Developed cross-departmental data analysis frameworks and reporting systems
- Created strategic recommendations based on complex dataset analysis
- Implemented data quality control measures and accuracy improvement initiatives
- Collaborated with stakeholders to optimize information system operations

ACHIEVEMENTS:-

Accenture Celebrates Excellence: GROWTH CATALYST AWARD

Participation in GM Skill Up, Scale Up AI Scalars Hackathon (Second position)

CERTIFICATIONS: -

- Career Essentials in Generative AI by Microsoft and LinkedIn
- Learning Microsoft 365 Copilot
- Ethics in the Age of Generative AI by LinkedIn
- Generative AI: The Evolution of Thoughtful Online Search
- AZ-900: Microsoft Azure Fundamentals Certification
- Databricks Partner Training Developer Essentials Badge (12/2021 - 12/2023)
- Databricks Partner Training Developer Foundations Badge (01/2022 - 01/2024)
- IBM- Data Analytics Certification

PERSONAL INFORMATION: -

Hobbies – Exploring Food, Fitness, Badminton

Language Known – English, Hindi, Punjabi

Date of Birth – 7th July 1998

Permanent Address – Maithri Layout, Whitefield, Bengaluru, 560066

